

AGENT CREATIVITY HACKATHON 2026 · WORKBUDDY TRACK

CaseCraft

Coordinated Support Packs from an Autism Passport —
one agent run, four audiences, one audited record.

Team CaseCraft · enochfyw@gmail.com · github.com/G-626/work_buddy_collab

Verified in the real WorkBuddy app on 2026-08-12 (Level-2). Demo data is fictional.

One child. Five adults. Five partial pictures.

An autism student is supported by parents, a class teacher, subject teachers, a social worker and a therapist. Each keeps **their own notes** — different names for the same trigger, different instructions for the same transition.

Reports arrive **late, duplicated and inconsistent**; a teacher gets a strategy memo, the parent gets a clinical phrase no one explained.

And the same sensitive detail can leak into the wrong audience's document — eroding trust, or worse.

The cost

Support depends on the **memory of whoever is in the room** — not on a shared, versioned record of what actually works for this child.

Demo data in this deck is fictional (Marco L., 14). No real child's information is shown anywhere.

The passport: one curated record, four coordinated views.

CaseCraft keeps a single **autism passport** per student — what supports them, their communication profile, triggers & sensory needs — curated by a human social worker, **versioned** (v1.0 → v1.1 → v1.2).

Every session adds a **passport delta**; every delivery renders four **audience-specific views** from the same source of truth:

Parent report — plain language · **Teacher guide** — classroom strategies · **Social story** — first person, for the student · **Therapist summary** — session evidence.

Observables, not verdicts

The machine extracts **what was observed** ("covered ears ×3 when the mixer started") and the **therapist interprets** ("builds tolerance with ear defenders").
No automated diagnosis, no microexpression-to-emotion claims.

Plan · Execute · Review · Deliver — one run, audited.

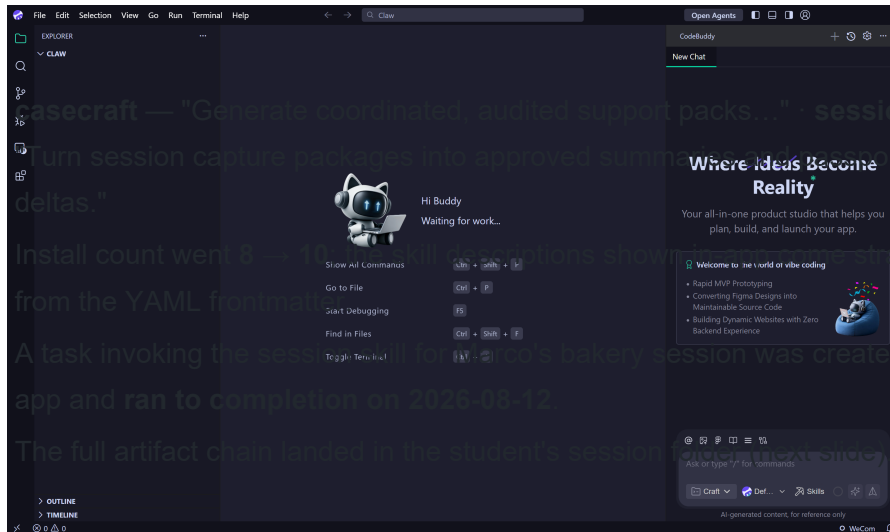
A WorkBuddy agent runs the **session skill** end to end from a capture package (transcript + observables):

- 1 · **Draft** — session summary with timestamps, counts, quotes, open questions for the worker.
- 2 · **Two-pass linter** — blocked-token scan + cross-audience parity (teacher & parent must say the same thing, never verbatim-identical script).
- 3 · **Review gate** — nothing merges without the social worker's sign-off (ADR-0004).
- 4 · **Passport delta** — v1.0 → v1.1 — then regenerate all four views + the linter report.

Human stays in the loop

The gate is **APPROVED with one edit** in the verified run — the machine proposes, the professional disposes.

Imported into the real WorkBuddy app — and it ran.



Both skills were imported into WorkBuddy v4.10.4 (logged in as enochfyw@gmail.com) via their native skill store.

WorkBuddy app — Skills tab, INSTALLED: 10 (casecraft, session highlighted).
Captured 2026-08-12.

What the agent produced — every file verified.

[students/marco/sessions/2026-08-08-bakery/](#)

01-session-draft.md — observed events w/ timestamps, counts, quotes, open questions.

02-review-gate.md — Decision: **APPROVED with one edit** (K. Wong, School Social Worker).

03-passport-v1.1.md — passport delta (1.0 → 1.1, Last updated 2026-08-12).

05-linter-report.md — **6/6 checks PASS**; leak drill **BLOCKED (FAIL)** as designed.

summary.json — outcome MET · checks 6/6 · drill BLOCKED.

views/ — parent guide, teacher guide, social story, therapist summary + linter report, rendered to PDF (34–43 KB each).

Draft content is agent-written (LLM variance); the linter + gate + engine make the output auditable regardless of that variance.

Deterministic quality gates against LLM variance.

A pure-LLM pipeline drifts. CaseCraft's linter is **code, not vibes** — it runs the same checks every time.

Blocked-token scan — clinical vocabulary ("assessment", "clinical", "goals file") is caught in every audience view.

Verbatim-script parity — teacher and parent guides must reference the same script but never share verbatim sentences.

Freshness — passport sections untouched >30 days are flagged for review.

The leak drill

A deliberately faulty parent draft was fed to the gate: **FAIL — delivery blocked** (tokens: assessment, clinical, goal progress, goals file). One blocked token blocks the whole pack until a human fixes it.

Then the pack was regenerated cleanly: **all views PASS** (see the 6/6 table).

Native skills, automation, and per-person delivery.

skill.yml manifests — each skill ships with the WorkBuddy-native metadata file (name, description, allowed-tools, config) plus SKILL.md — uploadable as a folder or .zip.

allowed-tools: Read, Write, Bash — the agent reads the session folder, writes the artifacts, and runs the engine scripts.

Automation tasks — a daily task checks each student's **next-session marker**; the day before a session it collects parent/teacher replies and drafts the pre-session brief (T-1 feedback loop).

Per-person delivery — each stakeholder receives only their view via their own messenger DM (Slack per-user DM = the per-login document, read-only by construction; replies feed the next hypothesis).

Secure by design — disclosed, not hidden.

Local transcription — audio is transcribed on-device (Whisper); the audio file never leaves the machine.

Remote summarisation, disclosed — WorkBuddy sends content to third-party LLMs (retention up to 14 days; Singapore/PRC transfer). We say "**secure by design**" — human-in-the-loop review + pseudonymised data + local transcription — never "everything stays local".

Observables only — no microexpression→emotion claims; the machine reports events, the therapist interprets.

Fictional demo data — Marco L. and Priya are fixtures; no real child's audio or identity is used.

Not in scope (stated, not hidden)

No hosted web app or logins — delivery is per-role files + messenger DMs. No automated diagnosis — the passport is a support record, not a clinical document.

From verified loop to pilot.

Capture — phone/tablet recording on the worker's device (the stop button lives there); whisper transcribes, the skill ingests.

Observables — MediaPipe landmarks → posture/gaze events in the capture package.

Delivery — Slack per-user DMs live; Automation feedback loop wired to next-session markers.

Pilot — one partner school, 3–5 students, 4 weeks: measure report turnaround (days → hours) and cross-audience consistency.

Today, verified

Level-1 engine (deterministic dry-run, 6/6 checks) **and** Level-2 in-app run — both green on 2026-08-12. Skills are uploadable to the portal as-is (folder or .zip with SKILL.md + YAML).

Team CaseCraft · enochfyw@gmail.com · github.com/G-626/work_buddy_collab