

AGENT CREATIVITY HACKATHON 2026 WORKBUDDY TRACK

CaseCraft

Coordinated support packs from an autism passport.

One child. Five adults. Five partial pictures.

There is a better way to understand them.

[VIEW THE REPO →](#)

VERIFIED IN THE REAL WORKBUDDY APP • 2026-08-13 • DEMO DATA IS FICTIONAL



Five adults. Five partial pictures.

Every student is surrounded by parents, a class teacher, subject teachers, a social worker and a therapist — each holding their own version of the same child.

Marco covers his ears when the mixer starts. That isn't misbehaviour — it's information.

- Different names for the same trigger. Different instructions for the same transition.
- Reports arrive late, duplicated and inconsistent — a teacher gets a strategy memo, a parent gets a clinical phrase nobody explained.
- The same sensitive detail can leak into the wrong audience's document — eroding trust, or worse.

Support depends on the memory of whoever is in the room — not on a shared, versioned record of what actually works for this child.



One passport. Four views.

A single curated autism passport per student — supports, communication profile, triggers and sensory needs — kept by a human social worker and versioned as it grows: v1.0 → v1.1 → v1.2.

PARENT REPORT

Plain language. What happened, what helped, what's next.

TEACHER GUIDE

Classroom strategies that work for this child.

SOCIAL STORY

First person, for the student to read.

THERAPIST SUMMARY

Session evidence, timestamps, counts, quotes.

Observables, not verdicts. The machine extracts what was observed — “covered ears ×3 when the mixer started”. The therapist interprets. No automated diagnosis. No microexpression-to-emotion claims.



Plan · Execute · Review · Deliver.

A WorkBuddy agent runs the casecraft skill end to end from a capture package — transcript plus observables. Both entry points (/session and /casecraft) live in the one skill.

01 · DRAFT

Session summary with timestamps, counts, quotes and open questions for the worker.

02 · TWO-PASS LINTER

Blocked-token scan + cross-audience parity. Teacher and parent must say the same thing — never verbatim-identical script.

03 · REVIEW GATE

Nothing merges without the social worker's sign-off. Human in the loop, always.

04 · PASSPORT DELTA

v1.0 → v1.1, then all four audience views regenerate in one run.

In the verified run, the gate came back APPROVED with one edit — the machine proposes, the professional disposes.



It ran — inside the real app.

● VERIFIED • 2026-08-13

WORKBUDDY V4.10.4

One combined skill imported through the app's native skill store. A task invoking /session for Marco's bakery session ran to completion, and the full artifact chain landed in the session folder.

```
$ skills install • INSTALLED → 9
casecraft One skill, whole loop — /session (capture→passport) + /casecraft (pack) commands

$ run /session → marco • 2026-08-08-bakery
draft • review-gate • passport-v1.1 • linter 6/6 • summary.json • views/ ... done
```



Every file verified.

students/marco/sessions/2026-08-08-bakery/ — what the agent produced, checked,
and rendered:

01-session-draft.md

observed events, timestamps, counts, quotes, open questions

02-review-gate.md

APPROVED with one edit — K. Wong, School Social Worker

03-passport-v1.1.md

passport delta — v1.0 → v1.1

05-linter-report.md

6/6 checks PASS · leak drill BLOCKED, as designed

summary.json

outcome MET · checks 6/6 · drill BLOCKED

views/

parent guide · teacher guide · social story · therapist summary — PDF,
34–43 KB each

The draft is agent-written — that variance is expected. The linter, gate and engine make the output auditable regardless.



Deterministic gates against LLM drift.

● LINTER 6/6 PASS

LEAK DRILL BLOCKED

A pure-LLM pipeline drifts. CaseCraft's linter is code, not vibes — the same checks, every run.

- Blocked-token scan — clinical vocabulary ("assessment", "clinical", "goals file") is caught in every audience view.
- Verbatim-script parity — teacher and parent guides must reference the same script, never share verbatim sentences.
- Freshness — passport sections untouched for over 30 days are flagged for review.

The drill: a deliberately faulty parent draft was fed to the gate — **FAIL**, delivery blocked (tokens: assessment, clinical, goal progress, goals file). One blocked token blocks the whole pack until a human fixes it.



Native skills. Automation. Per-person delivery.

SKILL.YML MANIFESTS

Each skill ships the WorkBuddy-native metadata file — name, description, allowed-tools, config — plus SKILL.md. Uploadable as a folder or a .zip.

AUTOMATION TASKS

A daily task checks each student's next-session marker. The day before a session it collects parent and teacher replies and drafts the pre-session brief.

No hosted web app, no custom logins. The per-user DM is the per-login document.

ALLOWED-TOOLS: READ, WRITE, BASH

The agent reads the session folder, writes the artifacts, and runs the engine scripts. A capability contract, not a daemon.

PER-PERSON DELIVERY

Each stakeholder receives only their view, via their own messenger DM — read-only by construction. Replies feed the next session's hypothesis.



Secure by design — disclosed, not hidden.

- Local transcription — audio is transcribed on-device (Whisper); the recording never leaves the worker's machine.
- Remote summarisation, disclosed — WorkBuddy sends content to third-party LLMs (retention up to 14 days, Singapore / PRC transfer). We say “secure by design” — human-in-the-loop review, pseudonymised data, local transcription. We never say “everything stays local”.
- Observables only — no microexpression-to-emotion claims. The machine reports events; the therapist interprets.
- Fictional demo data — Marco and Priya are fixtures. No real child's audio or identity anywhere.

Not in scope, stated plainly: no hosted web app or logins — delivery is per-role files plus messenger DMs. No automated diagnosis — the passport is a support record, not a clinical document.



From verified loop to pilot.

CAPTURE

Phone or tablet recording on the worker's device; whisper transcribes, the skill ingests.

OBSERVABLES

MediaPipe landmarks → posture and gaze events in the capture package.

DELIVERY

Messenger per-user DMs live; the automation feedback loop wired to next-session markers.

PILOT

One partner school, 3–5 students, 4 weeks — measure report turnaround from days to hours, and cross-audience consistency.

Today, verified: the Level-1 engine (deterministic dry-run; 6/6 checks) and the Level-2 in-app run — both green on 2026-08-13. The skills are uploadable to the portal as-is.

[START THE PILOT →](#)[VIEW THE REPO →](#)

Team CaseCraft · github.com/G-626/work_buddy_collab